

Ejercicio 1:

```
vagrant@oom2425-master:~ $ cat /etc/ansible/hosts
[masters]
oom2425-master
[workers]
oom2425-worker1
oom2425-worker2
# Group 'cluster' with all nodes
[cluster:children]
masters
workers
vagrant@oom2425-master:~ $ ansible cluster -m shell -a "hostname && uname -a"
oom2425-worker1 | CHANGED | rc=0 >>
oom2425-worker1
Linux oom2425-worker1 6.1.0-29-amd64 #1 SMP PREEMPT_DYNAMIC Debian 6.1.123-1 (2025-01-02) x86_64 GNU/Linux
oom2425-worker1 | CHANGED | rc=0 >>
oom2425-worker1
Linux oom2425-worker1 6.1.0-29-amd64 #1 SMP PREEMPT_DYNAMIC Debian 6.1.123-1 (2025-01-02) x86_64 GNU/Linux
oom2425-worker2 | CHANGED | rc=0 >>
oom2425-worker2
Linux oom2425-worker2 6.1.0-29-amd64 #1 SMP PREEMPT_DYNAMIC Debian 6.1.123-1 (2025-01-02) x86_64 GNU/Linux
vagrant@oom2425-master:~ $ ansible cluster -m shell -a "df -h | grep /data"
oom2425-worker1 | CHANGED | rc=0 >>
/dev/sdb              20G   24K   19G   1% /data/disk0
/dev/sdc              20G   24K   19G   1% /data/disk1
oom2425-worker2 | CHANGED | rc=0 >>
/dev/sdb              20G   24K   19G   1% /data/disk0
/dev/sdc              20G   24K   19G   1% /data/disk1
oom2425-master | CHANGED | rc=0 >>
/dev/sdb              20G   24K   19G   1% /data/disk0
/dev/sdc              20G   24K   19G   1% /data/disk1
vagrant@oom2425-master:~ $
```

Ejercicio 2:

```
vagrant@oom2425-master:~ $ sudo exportfs
/share/nfs          10.10.1.10/24
vagrant@oom2425-master:~ $
vagrant@oom2425-master:~ $ ansible workers -a "cat /proc/mounts" | grep "share/nfs"
oom2425-master:/share/nfs /share/nfs nfs4 rw,sync,relatime,vers=4.2,rsize=262144,wsize=262144,namlen=255,hard,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=10.10.1.11,local_lock=none,addr=10.10.1.10 0 0
oom2425-master:/share/nfs /share/nfs nfs4 rw,sync,relatime,vers=4.2,rsize=262144,wsize=262144,namlen=255,hard,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=10.10.1.12,local_lock=none,addr=10.10.1.10 0 0
vagrant@oom2425-master:~ $
```

Ejercicio 3:

```
vagrant@oom2425-master:~ $ su -l hadoop
Password:
hadoop@oom2425-master:~ $
hadoop@oom2425-master:~ $ echo $PATH && echo $HADOOP_HOME
/usr/local/bin:/usr/bin:/bin:/usr/local/games:/usr/games:/share/nfs/hadoop/hadoop-3.4.1/bin:/share/nfs/hadoop/hadoop-3.4.1/sbin
/share/nfs/hadoop/hadoop-3.4.1
hadoop@oom2425-master:~ $
hadoop@oom2425-master:~ $ tail -6 $HADOOP_HOME/etc/hadoop/hadoop-env.sh
# BEGIN ANSIBLE MANAGED BLOCK
export HDFS_NAMENODE_OPTS="-Xmx512m $HDFS_NAMENODE_OPTS"
export HDFS_SECONDARYNAMENODE_OPTS="-Xmx512m $HDFS_SECONDARYNAMENODE_OPTS"
export HDFS_DATANODE_OPTS="-Xmx512m $HDFS_DATANODE_OPTS"
export JAVA_HOME=/usr/lib/jvm/temurin-11-jdk-amd64
# END ANSIBLE MANAGED BLOCK
hadoop@oom2425-master:~ $ hadoop version
Hadoop 3.4.1
Source code repository https://github.com/apache/hadoop.git -r 4d7825309348956336b8f06a08322b78422849b1
Compiled by mthakur on 2024-10-09T14:57Z
Compiled on platform linux-x86_64
Compiled with protoc 3.23.4
From source with checksum 7292fe9dba5e2e44e3a9f763fce3e680
This command was run using /share/nfs/hadoop/hadoop-3.4.1/share/hadoop/common/hadoop-common-3.4.1.jar
hadoop@oom2425-master:~ $
```

Ejercicio 4:

```
hadoop@oom2425-master:~$ start-dfs.sh
Starting namenodes on [oom2425-master]
Starting datanodes
Starting secondary namenodes [oom2425-master]
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ ansible cluster -a "jps"
oom2425-worker1 | CHANGED | rc=0 >>
1939 Jps
1764 NodeManager
1652 DataNode
oom2425-worker2 | CHANGED | rc=0 >>
1933 Jps
1759 NodeManager
1647 DataNode
oom2425-master | CHANGED | rc=0 >>
3652 Jps
2936 NameNode
3083 SecondaryNameNode
3293 ResourceManager
hadoop@oom2425-master:~$
```

Ejercicio 5:

```
hadoop@oom2425-master:~$ ls -lh $HOME
total 715M
-rw-r--r-- 1 hadoop hadoop 715M Mar 25 11:34 wiki.xml
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ hdfs dfs -ls -R -h /
drwxr-xr-x  - hadoop supergroup          0 2025-05-06 17:32 /input
-rw-r--r--  2 hadoop supergroup    714.8 M 2025-05-06 17:27 /input/wiki.xml
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ ansible workers -m shell -a "df -h | grep /data"
oom2425-worker2 | CHANGED | rc=0 >>
/dev/sdb           20G  366M   19G   2% /data/disk0
/dev/sdc           20G  355M   19G   2% /data/disk1
oom2425-worker1 | CHANGED | rc=0 >>
/dev/sdc           20G  355M   19G   2% /data/disk1
/dev/sdb           20G  366M   19G   2% /data/disk0
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ hdfs dfs -df -h
Filesystem      Size   Used  Available  Use%
hdfs://oom2425-master:9000  78.1 G  1.4 G    72.6 G   2%
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$
```

Tema: Prácticas (50%) | GEI-AISI X AISI - Prácticas Hadoop - aisi24 X Namenode information X Apache Hadoop 3.4.1 – Overview X +

http://localhost:9870/dfshealth.html#tab-datanode 80% ☆

Datanode Information

In service Down Decommissioning Decommissioned Decommissioned & dead
Entering Maintenance In Maintenance In Maintenance & dead

Datanode usage histogram

Disk usage of each DataNode (%)

In operation

DataNode State All Show 25 entries Search:

Node	Http Address	Last contact	Last Block Report	Used	Non DFS Used	Capacity	Blocks	Block pool used	Block pool usage StdDev	Version
/default-rack/oom2425-worker1:9866 (10.10.1.11:9866)	http://oom2425-worker1:9864	2s	27m	720.56 MB	120 KB	39.04 GB	23	720.56 MB (1.8%)	0.03%	3.4.1
/default-rack/oom2425-worker2:9866 (10.10.1.12:9866)	http://oom2425-worker2:9864	1s	27m	720.56 MB	120 KB	39.04 GB	23	720.56 MB (1.8%)	0.03%	3.4.1

Showing 1 to 2 of 2 entries Previous 1 Next

```

hadoop@oom2425-master: ~ × + ▾

Job Counters
    Launched map tasks=23
    Launched reduce tasks=1
    Data-local map tasks=23
    Total time spent by all maps in occupied slots (ms)=6615360
    Total time spent by all reduces in occupied slots (ms)=1253736
    Total time spent by all map tasks (ms)=826920
    Total time spent by all reduce tasks (ms)=156717
    Total vcore-milliseconds taken by all map tasks=826920
    Total vcore-milliseconds taken by all reduce tasks=156717
    Total megabyte-milliseconds taken by all map tasks=423383040
    Total megabyte-milliseconds taken by all reduce tasks=80239104

Map-Reduce Framework
    Map input records=10863913
    Map output records=78462178
    Map output bytes=1037122162
    Map output materialized bytes=360463172
    Input split bytes=2438
    Combine input records=78462178
    Combine output records=14550927
    Reduce input groups=7941149
    Reduce shuffle bytes=360463172
    Reduce input records=14550927
    Reduce output records=7941149
    Spilled Records=43409223
    Shuffled Maps =23
    Failed Shuffles=0
    Merged Map outputs=23
    GC time elapsed (ms)=8661
    CPU time spent (ms)=393860
    Physical memory (bytes) snapshot=7183773696
    Virtual memory (bytes) snapshot=55225065472
    Total committed heap usage (bytes)=4305453056
    Peak Map Physical memory (bytes)=334389248
    Peak Map Virtual memory (bytes)=2312450048
    Peak Reduce Physical memory (bytes)=458141696
    Peak Reduce Virtual memory (bytes)=2314899456

Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0

File Input Format Counters
    Bytes Read=749627124
File Output Format Counters
    Bytes Written=213940419

real    4m18.408s
user    0m11.668s
sys     0m0.767s
hadoop@oom2425-master:~$
```

```

hadoop@oom2425-master: ~ × + ▾

hadoop@oom2425-master:~$ hadoop fs -ls -R -h /hadoop-outputR1
-rw-r--r--  2 hadoop supergroup          0 2025-05-06 17:56 /hadoop-outputR1/_SUCCESS
-rw-r--r--  2 hadoop supergroup  204.0 M 2025-05-06 17:56 /hadoop-outputR1/part-r-00000
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ hadoop fs -ls -R -h /hadoop-outputR4
-rw-r--r--  2 hadoop supergroup          0 2025-05-06 18:04 /hadoop-outputR4/_SUCCESS
-rw-r--r--  2 hadoop supergroup   50.9 M 2025-05-06 18:04 /hadoop-outputR4/part-r-00000
-rw-r--r--  2 hadoop supergroup   50.9 M 2025-05-06 18:04 /hadoop-outputR4/part-r-00001
-rw-r--r--  2 hadoop supergroup   51.1 M 2025-05-06 18:04 /hadoop-outputR4/part-r-00002
-rw-r--r--  2 hadoop supergroup   51.1 M 2025-05-06 18:04 /hadoop-outputR4/part-r-00003
hadoop@oom2425-master:~$
```

```

hadoop@oom2425-master: ~ × + ▾

hadoop@oom2425-master:~$ hadoop fs -getmerge -skip-empty-file /hadoop-outputR4 /tmp/hadoopOutputR4
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ ls -lh /tmp/hadoopOutputR*
-rw-r--r-- 1 hadoop hadoop 205M May  6 18:11 /tmp/hadoopOutputR1
-rw-r--r-- 1 hadoop hadoop 205M May  6 18:16 /tmp/hadoopOutputR4
hadoop@oom2425-master:~$ cat /tmp/hadoopOutputR1| grep "<page>" 
<page> 167269
hadoop@oom2425-master:~$ cat /tmp/hadoopOutputR4| grep "<page>" 
<page> 167269
hadoop@oom2425-master:~$
```

Ejercicio 6:

Ejercicio 7:

```
25/05/06 19:01:59 INFO SparkContext: Created broadcast 2 from broadcast at DAGScheduler.scala:1585
25/05/06 19:01:59 INFO DAGScheduler: Submitting 4 missing tasks from ResultStage 1 (MapPartitionsRDD[5] at map at IOCommon.scala:37) (first 15 tasks are for partitions Ve
ctor(0, 1, 2, 3))
25/05/06 19:01:59 INFO YarnScheduler: Adding task set 1.0 with 4 tasks resource profile 0
25/05/06 19:01:59 INFO TaskSetManager: Starting task 0.0 in stage 1.0 (TID 23) (oom2425-worker2, executor 2, partition 0, NODE_LOCAL, 9003 bytes)
25/05/06 19:01:59 INFO TaskSetManager: Starting task 1.0 in stage 1.0 (TID 24) (oom2425-worker1, executor 1, partition 1, NODE_LOCAL, 9003 bytes)
25/05/06 19:01:59 INFO TaskSetManager: Starting task 2.0 in stage 1.0 (TID 25) (oom2425-worker2, executor 2, partition 2, NODE_LOCAL, 9003 bytes)
25/05/06 19:01:59 INFO TaskSetManager: Starting task 3.0 in stage 1.0 (TID 26) (oom2425-worker1, executor 1, partition 3, NODE_LOCAL, 9003 bytes)
25/05/06 19:01:59 INFO BlockManagerInfo: Added broadcast_2_piece0 in memory on oom2425-worker2:39625 (size: 38.4 KiB, free: 148.3 MiB)
25/05/06 19:01:59 INFO BlockManagerInfo: Added broadcast_2_piece0 in memory on oom2425-worker1:38793 (size: 38.4 KiB, free: 148.3 MiB)
25/05/06 19:01:59 INFO MapOutputTrackerMasterEndpoint: Asked to send map output locations for shuffle 0 to 10.10.1.12:54302
25/05/06 19:01:59 INFO MapOutputTrackerMasterEndpoint: Asked to send map output locations for shuffle 0 to 10.10.1.11:50746
25/05/06 19:02:32 INFO TaskSetManager: Finished task 2.0 in stage 1.0 (TID 25) in 32227 ms on oom2425-worker2 (executor 2) (1/4)
25/05/06 19:02:32 INFO TaskSetManager: Finished task 0.0 in stage 1.0 (TID 23) in 33538 ms on oom2425-worker2 (executor 2) (2/4)
25/05/06 19:02:34 INFO TaskSetManager: Finished task 3.0 in stage 1.0 (TID 26) in 34844 ms on oom2425-worker1 (executor 1) (3/4)
25/05/06 19:02:34 INFO TaskSetManager: Finished task 1.0 in stage 1.0 (TID 24) in 35287 ms on oom2425-worker1 (executor 1) (4/4)
25/05/06 19:02:34 INFO DAGScheduler: ResultStage 1 (runJob at SparkHadoopWriter.scala:83) finished in 35.479 s
25/05/06 19:02:34 INFO YarnScheduler: Removed TaskSet 1.0, whose tasks have all completed, from pool
25/05/06 19:02:34 INFO DAGScheduler: Job 0 is finished. Cancelling potential speculative or zombie tasks for this job
25/05/06 19:02:34 INFO YarnScheduler: Killing all running tasks in stage 1: Stage finished
25/05/06 19:02:34 INFO DAGScheduler: Job 0 finished: runJob at SparkHadoopWriter.scala:83, took 95.672989 s
25/05/06 19:02:34 INFO SparkHadoopWriter: Start to commit write Job job_202505061900587157168554155681238_0005.
25/05/06 19:02:34 INFO SparkHadoopWriter: Write Job job_202505061900587157168554155681238_0005 committed. Elapsed time: 223 ms.
25/05/06 19:02:34 INFO SparkContext: Invoking stop() from shutdown hook
25/05/06 19:02:34 INFO SparkContext: SparkContext is stopping with exitCode 0.
25/05/06 19:02:34 INFO SparkUI: Stopped Spark web UI at http://oom2425-master:4040
25/05/06 19:02:34 INFO YarnClientSchedulerBackend: Interrupting monitor thread
25/05/06 19:02:34 INFO YarnClientschedulerBackend: Shutting down all executors
25/05/06 19:02:34 INFO YarnSchedulerBackend$YarnDriverEndpoint: Asking each executor to shut down
25/05/06 19:02:34 INFO YarnClientSchedulerBackend: YARN client scheduler backend Stopped
25/05/06 19:02:34 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
25/05/06 19:02:35 INFO MemoryStore: MemoryStore cleared
25/05/06 19:02:35 INFO BlockManager: BlockManager stopped
25/05/06 19:02:35 INFO BlockManagerMaster: BlockManagerMaster stopped
25/05/06 19:02:35 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
25/05/06 19:02:35 INFO SparkContext: Successfully stopped SparkContext
25/05/06 19:02:35 INFO ShutdownHookManager: Shutdown hook called
25/05/06 19:02:35 INFO ShutdownHookManager: Deleting directory /tmp/spark-facc629a-3463-4452-9ffc-19e83a926377
25/05/06 19:02:35 INFO ShutdownHookManager: Deleting directory /tmp/spark-6e746948-9378-4db9-993d-a3c21f8628bf

real    2m17.623s
user    0m23.763s
sys     0m3.999s
hadoop@oom2425-master:~$
```

```

hadoop@oom2425-master:~$ hadoop fs -ls -R -h /spark-outputR1
-rw-r--r-- 2 hadoop supergroup 0 2025-05-06 18:56 /spark-outputR1/_SUCCESS
-rw-r--r-- 2 hadoop supergroup 204.0 M 2025-05-06 18:56 /spark-outputR1/part-r-00000
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ hadoop fs -ls -R -h /spark-outputR4
-rw-r--r-- 2 hadoop supergroup 0 2025-05-06 19:02 /spark-outputR4/_SUCCESS
-rw-r--r-- 2 hadoop supergroup 51.1 M 2025-05-06 19:02 /spark-outputR4/part-r-00000
-rw-r--r-- 2 hadoop supergroup 51.1 M 2025-05-06 19:02 /spark-outputR4/part-r-00001
-rw-r--r-- 2 hadoop supergroup 50.8 M 2025-05-06 19:02 /spark-outputR4/part-r-00002
-rw-r--r-- 2 hadoop supergroup 50.9 M 2025-05-06 19:02 /spark-outputR4/part-r-00003
hadoop@oom2425-master:~$ 

```

```

hadoop@oom2425-master:~$ hadoop fs -getmerge /spark-outputR4 /tmp/sparkOutputR4
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ ls -lh /tmp/*OutputR*
-rw-r--r-- 1 hadoop hadoop 205M May 6 18:11 /tmp/hadoopOutputR1
-rw-r--r-- 1 hadoop hadoop 205M May 6 18:16 /tmp/hadoopOutputR4
-rw-r--r-- 1 hadoop hadoop 205M May 6 19:13 /tmp/sparkOutputR4
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ cat /tmp/sparkOutputR4 | grep "<page>" 
<page> 167269
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ 

```

Ejercicio 8:

```

hadoop@oom2425-master:~$ hdfs dfs -ls -R -h /quijote-input
-rw-r--r-- 2 hadoop supergroup 1.0 M 2025-05-06 19:20 /quijote-input/el_quijote.txt
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ hdfs dfs -ls -R -h /quijote-output
-rw-r--r-- 2 hadoop supergroup 0 2025-05-06 19:25 /quijote-output/_SUCCESS
-rw-r--r-- 2 hadoop supergroup 251.2 K 2025-05-06 19:25 /quijote-output/part-r-00000
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ hdfs dfs -get /quijote-output $HOME
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ ls -lh $HOME
total 747M
-rw-r--r-- 1 hadoop hadoop 1.1M Mar 28 11:24 el_quijote.txt
drwxr-xr-x 2 hadoop hadoop 4.0K May 6 19:27 quijote-output
-rw-r--r-- 1 hadoop hadoop 31M Feb 7 14:39 sparkbench.jar
-rw-r--r-- 1 hadoop hadoop 715M Mar 25 11:34 wiki.xml
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ 
hadoop@oom2425-master:~$ sort -k2 -nr $HOME/quijote-output/part-r-00000 | head -n 5
que      10351
de       8947
y        8042
la       4941
a        4725
hadoop@oom2425-master:~$ cat $HOME/quijote-output/part-r-00000 | grep Rocinante
Rocinante    43
Rocinante,   67
Rocinante.   6
Rocinante;   3
hadoop@oom2425-master:~$ 

```

Ejercicio 9:

```
hadoop@oom2425-master:~ $ hadoop fs -rm -R /hadoop-output*
Deleted /hadoop-outputR1
Deleted /hadoop-outputR4
hadoop@oom2425-master:~ $ hadoop fs -rm -R /spark-output*
Deleted /spark-outputR1
Deleted /spark-outputR4
hadoop@oom2425-master:~ $ hadoop fs -rm -R /quijote-output*
Deleted /quijote-output
hadoop@oom2425-master:~ $ hadoop fs -df -h /
Filesystem           Size   Used  Available  Use%
hdfs://oom2425-master:9000  78.1 G  1.5 G   69.4 G   2%
hadoop@oom2425-master:~ $ hadoop fs -ls -R /
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 18:35 /input
-rw-r--r--  2 hadoop supergroup 31674614 2025-05-06 18:35 /input/sparkbench.jar
-rw-r--r--  2 hadoop supergroup 749528820 2025-05-06 17:27 /input/wiki.xml
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 19:20 /quijote-input
-rw-r--r--  2 hadoop supergroup 1060259 2025-05-06 19:20 /quijote-input/el_quijote.txt
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 17:52 /tmp
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 17:52 /tmp/hadoop-yarn
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 17:52 /tmp/hadoop-yarn/staging
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 17:52 /tmp/hadoop-yarn/staging/hadoop...
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 19:25 /tmp/hadoop-yarn/staging/hadoop...staging
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 17:52 /tmp/hadoop-yarn/staging/history
drwxrwxrwt - hadoop supergroup  0 2025-05-06 17:52 /tmp/hadoop-yarn/staging/history/done_intermediate
drwxrwx---  - hadoop supergroup  0 2025-05-06 19:25 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop
-rwrxrwx--- 2 hadoop supergroup 108857 2025-05-06 17:56 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0001-1746553967658-hadoop-word+co
unt-1746554213303-23-1-SUCCEEDED-root.default-1746553984594.jhist
-rwxrwx---  2 hadoop supergroup 447 2025-05-06 17:56 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0001.summary
-rwxrwx---  2 hadoop supergroup 306255 2025-05-06 17:56 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0001.conf.xml
-rwxrwx---  2 hadoop supergroup 121546 2025-05-06 18:04 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0002-1746554466828-hadoop-word+co
unt-1746554680081-23-4-SUCCEEDED-root.default-1746554477782.jhist
-rwxrwx---  2 hadoop supergroup 447 2025-05-06 17:52 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0002.summary
-rwxrwx---  2 hadoop supergroup 306255 2025-05-06 18:04 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0002.conf.xml
-rwxrwx---  2 hadoop supergroup 22942 2025-05-06 19:25 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0008-1746559530283-hadoop-word+co
unt-1746559536787-1-1-SUCCEEDED-root.default-1746559518689.jhist
-rwxrwx---  2 hadoop supergroup 441 2025-05-06 19:25 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0008.summary
-rwxrwx---  2 hadoop supergroup 306261 2025-05-06 19:25 /tmp/hadoop-yarn/staging/history/done_intermediate/hadoop/job_1746551854173_0008.conf.xml
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 18:40 /user
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 18:40 /user/hadoop
drwxr-xr-x  - hadoop supergroup  0 2025-05-06 19:02 /user/hadoop/.sparkStaging
hadoop@oom2425-master:~ $
```